

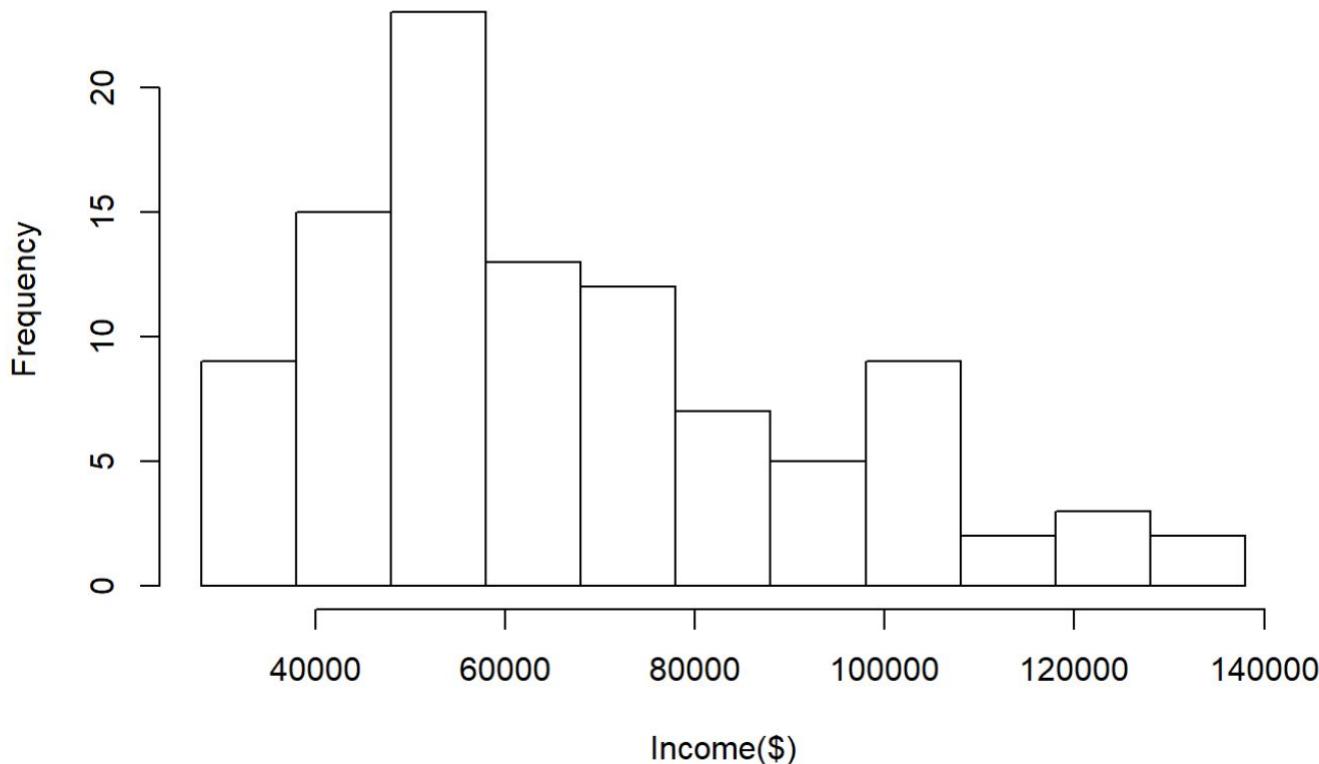
Exploratory Data Analysis

Faaz Arshad





Histogram of customers' annual salary



```
# firstly check the salary payment frequency of each customer
df_inc = data.frame(customer_id= unique(df_csmpl$customer_id)) #create a data frame to store result

# create a mode function that will be used to find out what is the salary payment frequency
Mode <- function(x) {
  ux <- unique(x)
  ux[which.max(tabulate(match(x, ux)))]
}

# Loop through all salary payment for each customer
# assume the salary Level is constant for each customer over the observed period
for (i in seq(nrow(df_inc))){
  trans_data <- df[df$customer_id == df_inc$customer_id[i]
                  & df$txn_description=='PAY/SALARY',c("amount","date")] %>%
    group_by(date) %>%
    summarise(amount = sum(amount))

  total_s <- sum(trans_data$amount)
  count = dim(trans_data)[1]
  if ( count == 0){
    df_inc$freq[i] = NA
    df_inc$level[i] = NA
  } else {
    s=c()
    lvl = c()
    for (j in seq(count-1)){
      s = c(s,(trans_data$date[j+1]-trans_data$date[j]))
      lvl = c(lvl,trans_data$amount[j])
    }
    lvl = c(lvl,tail(trans_data$amount,n=1))
    df_inc$freq[i] = Mode(s)
    df_inc$level[i] = Mode(lvl)
  }
}
```



Continued

```
df_inc$annual_salary= df_inc$level / df_inc$freq *365.25

# visualise the distribution of customers' annual salary
hist(df_inc$annual_salary[!is.na(df_inc$annual_salary)],breaks=c(seq(28000,140000,by = 10000)),
  main = "Histogram of customers' annual salary", xlab= 'Income($)')
```



Explore correlations between annual salary and various customer attributes

```
# create a dataframe to store relevant features for customers

df_cus <- df_csmp %>% # use df_csmp to summarize customers' consumption behavior
  select (customer_id,gender,age,amount,date,balance) %>%
  group_by(customer_id) %>%
  mutate(avg_no_weekly_trans= round(7*n()/length(unique(df$date)),0),max_amt = max(amount),
        no_large_trans = sum(amount>100), # an arbitrary $100 benchmark is selected based on the
        e
                                # transaction amount histogram created in task 1.3
  use_no_day=length(unique(date)),
  avg_trans_amt = mean(amount, na.rm =TRUE),
  med_bal = median(balance,na.rm=TRUE)) %>%
  select(-c("amount","date","balance")) %>%
  unique()

# create additional features
df_cus$age_below20 <- ifelse(df_cus$age<20,1,0)
df_cus$age_btwn20n40 <- ifelse(df_cus$age>=20 & df_cus$age <40,1,0)
df_cus$age_btwn40n60 <- ifelse(df_cus$age>=40 & df_cus$age <60,1,0)

# investigate the state where customers live
# assume they live where most transactions occurred (indicated by merchant_state)
df_region <- df_csmp %>%
  group_by(customer_id,merchant_state) %>%
  summarize(trans_count=n()) %>%
  group_by(customer_id) %>%
  mutate (no_state = n()) %>%
  filter(trans_count == max(trans_count))
```

```
# For equal number of transactions between multiple States, pick the most Likely State
n_occur = data.frame(table(df_region$customer_id))
cus_id_rep = n_occur$Var1[n_occur$Freq > 1]

state_by_cust_no <- rev(names(sort(table(df_region$merchant_state),rev = TRUE)))
t = data.frame(customer_id = cus_id_rep, merchant_state=NA)

for (i in seq(length(cus_id_rep))){
  s = df_region$merchant_state[df_region$customer_id == cus_id_rep[i]]
  for (state in state_by_cust_no){
    if (state %in% s){
      t[i,2] = state
      break
    }
  }
}

df_region <- df_region[!(df_region$customer_id %in% cus_id_rep), c(1,2)] %>%
  as.data.frame() %>%
  rbind(t) %>%
  rename( State = merchant_state)

# merge all the features into single dataframe
df_cus <- df_cus %>% merge(df_inc) %>%
```



Continued

```
merge(df_region)

# extract relevant features
df_cus_attr <- df_cus %>%
    select("gender", "annual_salary", "age", "avg_no_weekly_trans", "max_amt",
           "no_large_trans", "use_no_day", "avg_trans_amt", "med_bal", "State")
plot(df_cus_attr)
```



Decision Tree

Listed the code under instructions

```
# split into train and test datasets
smp_size <- floor(0.75 * nrow(df_cus))

## set the seed to make your partition reproducible
set.seed(123)
train_ind <- sample(seq_len(nrow(df_cus)), size = smp_size)

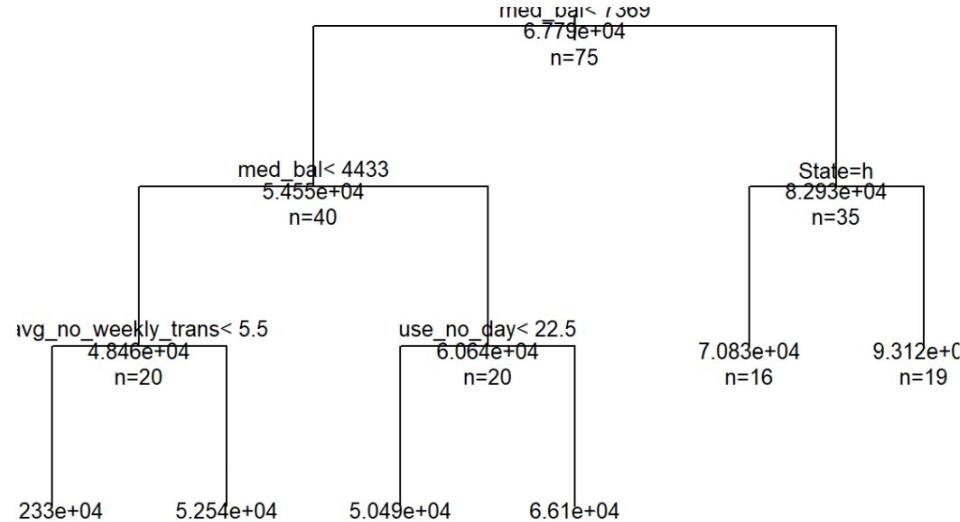
df_cus_train <- df_cus[train_ind, ]
df_cus_test <- df_cus[-train_ind, ]

fit3 <- rpart(annual_salary ~gender + age + avg_no_weekly_trans + max_amt + no_large_trans + us
e_no_day + avg_trans_amt + med_bal + age_below20 + age_btwn20n40 + age_btwn40n60 + State, method=
"anova", data=df_cus_train)

# plot tree
plot(fit3, uniform=TRUE,
      main="Regression Tree for Annual Salary ")
text(fit3, use.n=TRUE, all=TRUE, cex=.8)
```



Regression Tree for Annual Salary





```
# examine the prediction accuracy  
rmse(fit3, df_cus_test)
```

```
## [1] 25672.77
```



THE END